

TrustedRouter for enterprise

One connection to hundreds of models, bought on a public exchange, with proof that nobody read your prompts.


The problem with how you buy AI today

Every prompt your teams send is strategy, client data or source code. The intermediaries log it, and for the largest of them that traffic is the business model. At the same time your security review has approved two or three suppliers, so you queue behind everyone else for capacity and accept whatever the list price is. The shortage is trust, and the price is the symptom.

Why an exchange rather than another router

The routers and gateways competing for this decision all solve one narrow problem: software that chooses a model. You would still have to find supply, negotiate each contract and verify each provider yourself. That is a cost centre, and it is the same product from every vendor offering it.

An exchange inverts the relationship. Supply comes to you, at a public price, from many sellers competing for the order. Every unit is graded before it can be quoted: model, context, attestation class and jurisdiction. We verify that grade ourselves rather than accepting a provider's claim, and we have found encryption faults that suppliers then corrected.



Open source on both sides. Buyers connect through free harnesses and gateways; providers plug in through the free provider stack. One exchange clears the middle.

Security: open source and attestation together

Every competing router is closed source. You cannot read the code that touches your prompts, so privacy is an assertion in a contract. Publishing source alone does not fix this either, because nothing proves the code you read is the binary that ran.

We close the loop by combining both. The gateway, the provider stack and the attestation path are open source, so your engineers can read exactly what happens to a request. The running enclave then publishes its source commit, image digest and hardware measurement, so you can verify the deployed binary is built from the code you reviewed. Neither property is sufficient alone; together they turn privacy from a promise into something your security team confirms independently.

Fails closed. If attestation does not verify, the request does not execute. There is no silent fallback to an unverified path.

SOC 2 in progress. We are currently inside our SOC 2 Type II observation window, with the report to follow. DPA and BAA terms

Deployable inside your perimeter. Because the stack is open, you can run the whole thing in your own infrastructure rather than trusting our operations.

Controls on every request. Bring your own keys, set jurisdiction and provenance requirements, and require zero retention per call.

are available now.

Intelligence: the right model, not merely the cheapest

Most routers optimise on one axis, price, and send each request to whichever model is currently least expensive. For a consumer application that is acceptable. For a regulated workload it is a governance failure waiting to happen, because the cheapest capable model for a given task is frequently one your policy does not permit, hosted in a jurisdiction you have not approved, and a cost-only router will route there silently and without a record.

We route on fitness for the use case. Capability on the specific task, context length, latency, and then price. Provenance and jurisdiction are constraints you set per request rather than outcomes you discover afterwards. A workload can be restricted to a named set of providers, to models of a given origin, or to a single legal jurisdiction, and every request carries a record of which model ran where. Intelligence and control stop being a trade-off.

You can sell as well as buy

If you hold committed capacity you cannot use, or GPUs sitting idle, list them. Our open-source software runs the fleet from the hardware up, or we operate it for you. Capacity that would have expired earns a public bid instead, from buyers who never see what runs on it.

GETTING STARTED TAKES ONE LINE

We are OpenAI-compatible. Change **base_url**, keep your existing SDK and code. Start on a self-serve wallet from \$0, with migration credits for teams moving real volume. Enterprise terms, including a DPA and a BAA, are available on request.

References and further reading

Axios, *Routing is coming for the frontier AI labs*, 25 August 2026 · [axios.com/2026/08/25/routing-is-coming-for-the-frontier-ai-labs](#)
Axios, *AI startup TrustedRouter raises \$1.25 million*, 31 August 2026 · [axios.com/2026/08/31](#)
Open-source repositories and published attestation · [github.com/Lore-Hex](#) · [trustedrouter.com/docs](#)
Model listing and custom models · [trustedrouter.com/docs/custom-models](#)
On synthetic data derived from customer traffic: *Generative Data Refinement*, Google DeepMind, arXiv:2509.08653

FEATURED IN

AXIOS

BACKED BY

SLOW

Upfront Ventures · Avalanche VC · Learn Capital
Bill Tai · Linda Avey · engineers from OpenAI and Anthropic